

SPARK

@CrazyJvm

What is Spark

- a fast and general-purpose cluster computing system.
- high efficiency
- high level api (Scala, Java, Python)

How to run

- Local
- Standalone
- Mesos
- YARN

and an interactive shell (Scala supported)

About Scala

- JVM based
- Statically typed
- Interoperate with Java (vice-versa)

try in interactive shell !

The most important concept : RDD

RDDs : resilient distributed datasets

internally, each RDD is characterized by five main properties:

- * - A list of partitions**
- * - A function for computing each split**
- * - A list of dependencies on other RDDs**
- * - Optionally, a Partitioner for key-value RDDs (e.g. to say that the RDD is hash-partitioned)**
- * - Optionally, a list of preferred locations to compute each split on (e.g. block locations for an HDFS file)**

The most important concept : RDD

- immutable collections of objects spread across a cluster
- build through parallel transformations
- automatically rebuild on failure
- different storage level (memory management)


Over view

- RDDs
- Transformations (Lazy evaluation!!!)
- Action (def runJob[T, U: ClassManifest]
(rdd: RDD[T],
func: Iterator[T] => U): Array[U])

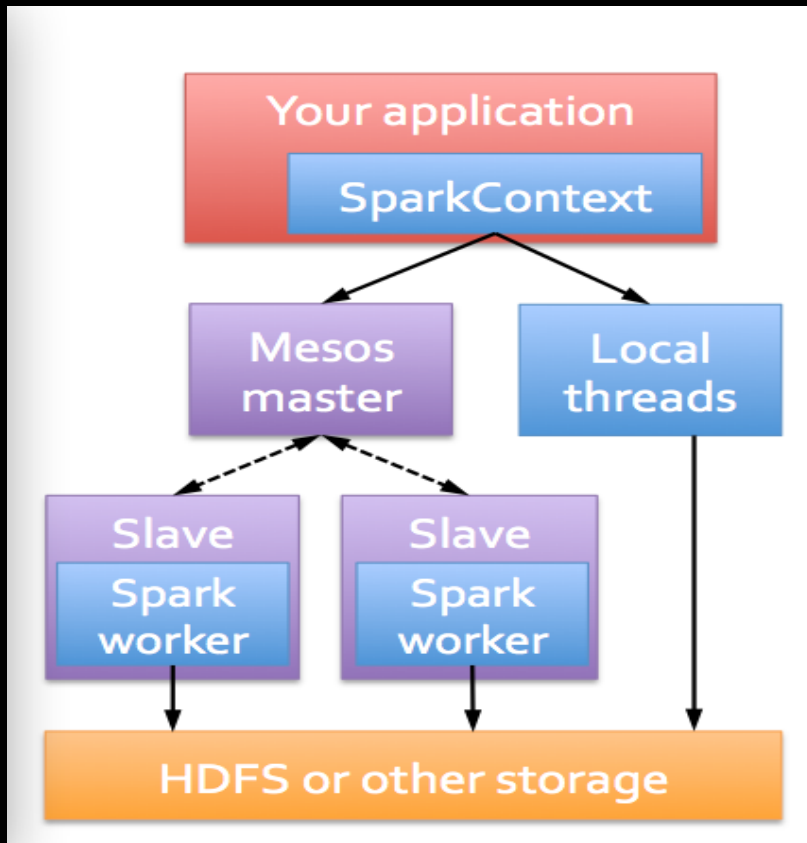
RDD:transformations & actions

Transformations	<div><div><i>map</i>(<i>f</i> : <i>T</i> ⇒ <i>U</i>) : RDD[<i>T</i>] ⇒ RDD[<i>U</i>]</div><div><i>filter</i>(<i>f</i> : <i>T</i> ⇒ Bool) : RDD[<i>T</i>] ⇒ RDD[<i>T</i>]</div><div><i>flatMap</i>(<i>f</i> : <i>T</i> ⇒ Seq[<i>U</i>]) : RDD[<i>T</i>] ⇒ RDD[<i>U</i>]</div><div><i>sample</i>(<i>fraction</i> : Float) : RDD[<i>T</i>] ⇒ RDD[<i>T</i>] (Deterministic sampling)</div><div><i>groupByKey</i>() : RDD[(<i>K</i>, <i>V</i>)] ⇒ RDD[(<i>K</i>, Seq[<i>V</i>])]</div><div><i>reduceByKey</i>(<i>f</i> : (<i>V</i>, <i>V</i>) ⇒ <i>V</i>) : RDD[(<i>K</i>, <i>V</i>)] ⇒ RDD[(<i>K</i>, <i>V</i>)]</div><div><i>union</i>() : (RDD[<i>T</i>], RDD[<i>T</i>]) ⇒ RDD[<i>T</i>]</div><div><i>join</i>() : (RDD[(<i>K</i>, <i>V</i>)], RDD[(<i>K</i>, <i>W</i>)]) ⇒ RDD[(<i>K</i>, (<i>V</i>, <i>W</i>))]</div><div><i>cogroup</i>() : (RDD[(<i>K</i>, <i>V</i>)], RDD[(<i>K</i>, <i>W</i>)]) ⇒ RDD[(<i>K</i>, (Seq[<i>V</i>], Seq[<i>W</i>]))]</div><div><i>crossProduct</i>() : (RDD[<i>T</i>], RDD[<i>U</i>]) ⇒ RDD[(<i>T</i>, <i>U</i>)]</div><div><i>mapValues</i>(<i>f</i> : <i>V</i> ⇒ <i>W</i>) : RDD[(<i>K</i>, <i>V</i>)] ⇒ RDD[(<i>K</i>, <i>W</i>)] (Preserves partitioning)</div><div><i>sort</i>(<i>c</i> : Comparator[<i>K</i>]) : RDD[(<i>K</i>, <i>V</i>)] ⇒ RDD[(<i>K</i>, <i>V</i>)]</div><div><i>partitionBy</i>(<i>p</i> : Partitioner[<i>K</i>]) : RDD[(<i>K</i>, <i>V</i>)] ⇒ RDD[(<i>K</i>, <i>V</i>)]</div></div>
Actions	<div><div><i>count</i>() : RDD[<i>T</i>] ⇒ Long</div><div><i>collect</i>() : RDD[<i>T</i>] ⇒ Seq[<i>T</i>]</div><div><i>reduce</i>(<i>f</i> : (<i>T</i>, <i>T</i>) ⇒ <i>T</i>) : RDD[<i>T</i>] ⇒ <i>T</i></div><div><i>lookup</i>(<i>k</i> : <i>K</i>) : RDD[(<i>K</i>, <i>V</i>)] ⇒ Seq[<i>V</i>] (On hash/range partitioned RDDs)</div><div><i>save</i>(<i>path</i> : String) : Outputs RDD to a storage system, <i>e.g.</i>, HDFS</div></div>

Spark runtime



Components



Just do it

- interactive shell
- local mode (get local data)
- standalone mode (get data from hdfs)
- programming in IDE(eclipse,idea)

Word Count

```
val text = sc.textFile("README.md")  
val wc = text.flatMap(_.split(" ")).map((_, 1)).reduceByKey(_+_)  
wc.collect
```

notice: reduceByKey is called by implicit conversion

```
implicit def rddToPairRDDFunctions[K: ClassManifest, V: ClassManifest]  
  (rdd: RDD[(K, V)]) = new PairRDDFunctions(rdd)
```

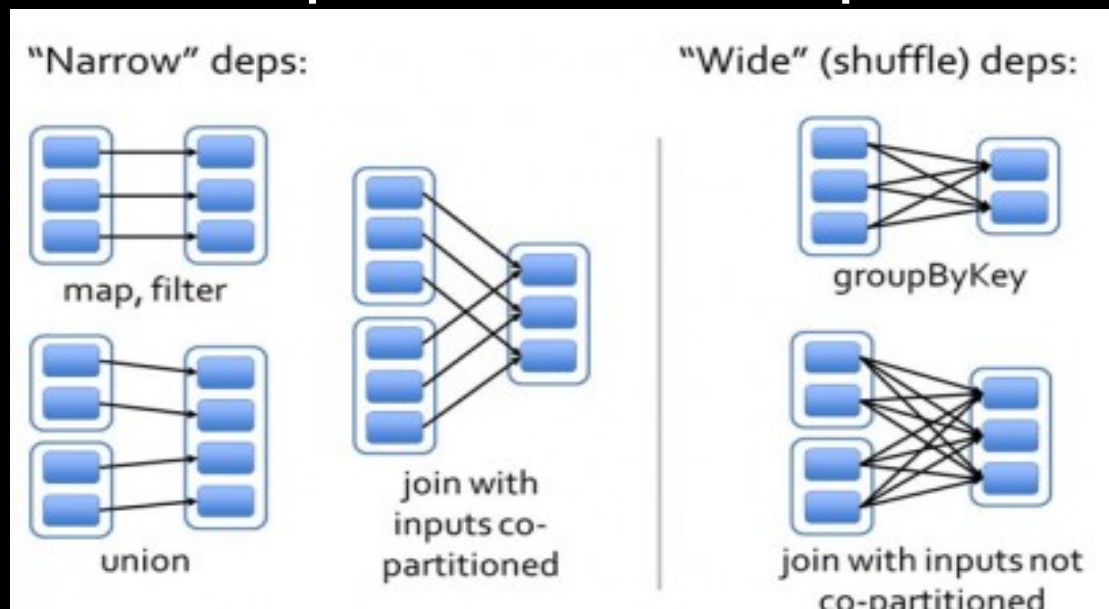
if we let wc.cache, what will happen?

RDD Lineage

- Narrow dependency: each partition of the parent RDD is used by at most one partition of the child RDD.
- Wide dependency: multiple child partitions may depend on a partition of parent RDD.

RDD Lineage

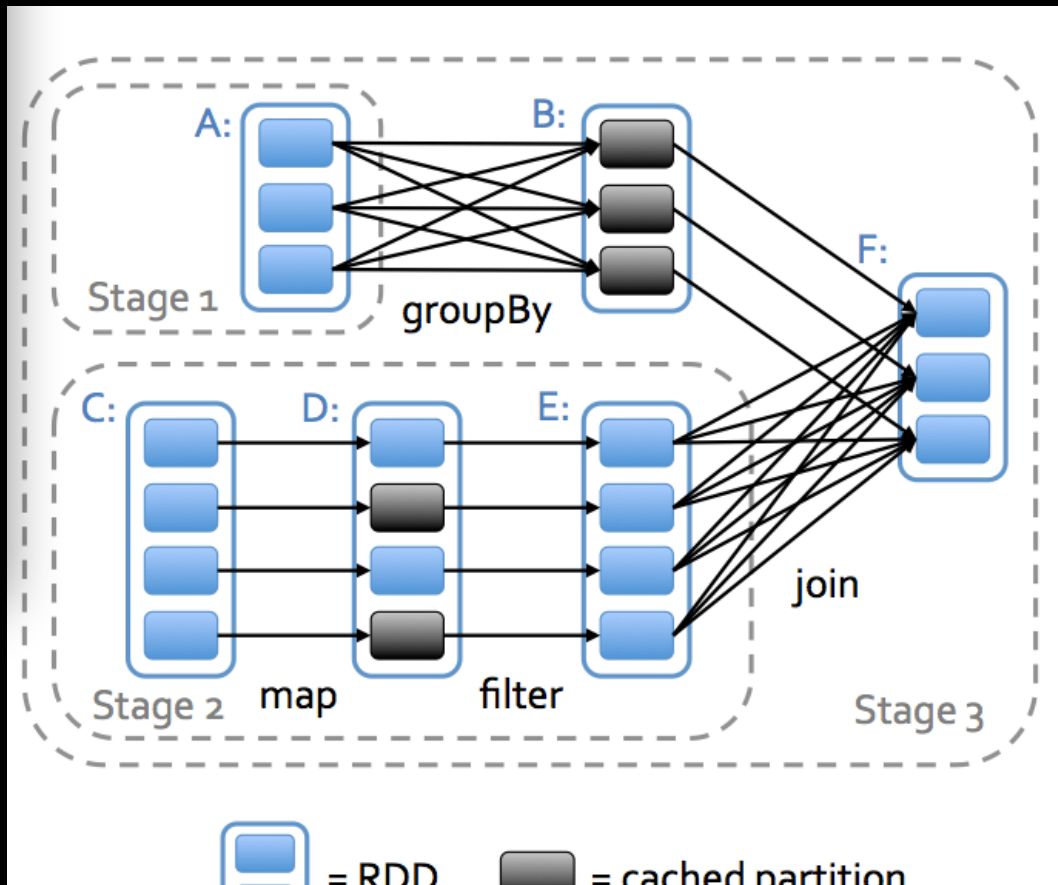
- optimization? -> pipeline
- the importance of co-partitioned



Task scheduler

- run general task graphs
- pipeline functions where possible
- Cache-aware data reuse and locality
- Partitioning-aware to avoid shuffles

Task scheduler



Schedule process

- RDD objects
- DAGScheduler
- TaskScheduler
- Worker

RDD fault tolerance

- recovery by lineage
- checkpoint

Q & A

thanks!